



(a) The winner depends on the model



(b) Output error vs. task score
within-cell ρ : -0.16 (12 methods); -0.38 (8 token-selective)



(c) How answers fail (Llama, 128k)

