



(a) The winner depends on the model



(b) Output error vs. task score
within-cell ρ : -0.05 (8 methods); 8k run: -0.36 (12)



eviction quantization SIEVE (mixed)

(c) How answers fail (Llama, 128k)

