

EVICT, QUANTIZE, OR ALLOCATE?

A PHASE DIAGRAM FOR KV-CACHE COMPRESSION

Anonymous authors

Paper under double-blind review

ABSTRACT

KV-cache compression has split into two literatures: *quantization* keeps every token at reduced precision, *eviction* keeps few tokens at full precision. We show they are two corners of one rate-allocation problem. Posed against the distortion of the attention *output* rather than of the keys, the optimum is reverse water-filling, $b_i^* = \log_2 a_i + c$, and eviction is its zero-rate tier at a cost $c_0=1$ that is derived rather than tuned. The same cost comparison yields a regime variable: the fraction of heads whose 2-bit tier is already costlier than eviction ($\tau^2 c_2 > c_0$). Across 10,240 heads of six LLMs at 8k–128k context, measured by exact recomputation, this *dead-tier fraction* orders the gain of allocation over the best corner at Spearman $\rho = -0.96$ (-1.00 within every model), contains no context length, and correctly anticipated a held-out collapse at 128k. Context length turns out to be a phase variable: a single model crosses from the allocation-dominated to the eviction-dominated regime as L grows, so “method X beats Y on model M” is ill-posed without a stated L . We turn the map into SIEVE, a per-head router calibrated in one forward pass. **TBD [R8]: end-task result, one sentence.**

1 INTRODUCTION

At long context the KV cache, not the weights, bounds batch size and decode throughput. Two families of methods compete to shrink it. Quantizers (Liu et al., 2024; Hooper et al., 2024; Zandieh et al., 2024; 2025) store every token at 2–4 bits; evictors (Zhang et al., 2023; Liu et al., 2023; Li et al., 2024; Xiao et al., 2024) store a subset of tokens at full precision and drop the rest. Each reports wins over the other, and practitioners pick by leaderboard.

These are not competing ideas; they are the endpoints of one. Attention does not consume keys, it consumes $o = \sum_i a_i v_i$. Once the objective is the error in o , the question “how many bits should token i receive?” has a closed-form answer that ranges over $\{0, \dots, b_{\max}\}$. Uniform quantization is the answer when every token gets the same width; eviction is the answer when only 0 and $b_{\max}$ survive. Neither family is right or wrong in general. The useful question is *when* each one is optimal, and whether that can be read off the model before deploying anything.

Our answer is a phase diagram (Fig. 1). The optimal bit profile has three shapes (Fig. 1a): flat when attention is diffuse, a staircase (the *productive band*, where mixed precision beats both corners), and keep/evict when attention is sharp. The boundary between the last two is not an empirical threshold. It follows from the same inequality that produces the allocation: a precision tier is worth buying only while its logit noise $\tau^2 c_b$ costs less than evicting the token ($c_0 = 1$). The fraction of heads for which the 2-bit tier fails this test is our order parameter, and Fig. 1b shows 24 measured configurations collapsing onto it.

Contributions.

1. **One objective for both literatures** (§2). Under output distortion, eviction is the zero-rate tier of a water-filling allocator with a derived cost $c_0 = 1$, and the optimal rule $b_i^* = \log_2 a_i + c$ reads only attention weights.
2. **A derived order parameter** (§3). The dead-tier fraction predicts where allocation pays across six architectures and five octaves of context ($\rho = -0.96$ pooled, -1.00 within each model). It contains no L , holds when the eviction baseline is redefined, and passed a held-out test.
3. **Context length is a phase variable.** τ rises monotonically with L in every model. Rising τ kills precision tiers from the bottom, so one model crosses regimes as context grows and a per-model verdict is ill-posed.
4. **SIEVE, a router built from the map** (§4). One calibration pass assigns each head to uniform, ladder, or eviction. It rests on a cost model that forward-predicts measured gains to within 0.79 – $1.34\times$, and on measured design facts (Table 1).




Figure 1: **(a)** Optimal bit-widths from reverse water-filling at a 3-bit average budget (synthetic Gaussian logits, measured cost constants). As logit spread τ grows the profile goes from near-uniform to a ladder to keep/evict. **(b)** *Measured*: routed gain of allocation over the better of uniform quantization and eviction, per (model, context); marker size grows with context and arrows mark the last doubling. One derived variable orders all 24 points ($\rho = -0.96$), and every model moves right as context grows. **(c)** *Projected, not measured (TBD [R8])*: the end-task picture the phase diagram predicts, with the largest gain at short-to-mid context, shrinking toward eviction at 128k.

2 ONE OBJECTIVE: EVICTION IS QUANTIZATION AT ZERO RATE

Output distortion. For query q and keys k_i , let $s_i = q^\top k_i / \sqrt{d}$, $a = \text{softmax}(s)$, $o = \sum_i a_i v_i$. Storing k_i at b_i bits perturbs the logit by δ_i with $\text{Var}(\delta_i) = \tau^2 c_{b_i}$. Here τ is the head’s logit spread and c_b the quantizer’s relative distortion; we use the variance because softmax is shift-invariant, so bias costs nothing. Since $\partial o / \partial s_i = a_i(v_i - o)$, to first order

$$\mathbb{E} \|\Delta o\|^2 \approx \sum_i a_i^2 \|v_i - o\|^2 \tau^2 c_{b_i}. \quad (1)$$

Evicting a set E and renormalizing costs $\sum_{i \in E} a_i^2 \|v_i - o\|^2$ to the same order. That is Eq. 1 with $\tau^2 c_0 = 1$, so eviction is a bit-width like any other and **its cost constant is derived, not chosen**.

Reverse water-filling. Minimizing Eq. 1 subject to $\sum_i b_i \leq BL$ is a classical rate-allocation problem (Cover & Thomas, 2006). The Lagrangian separates per token, and with $c_b \propto 4^{-b}$ its solution above the water level is $b_i^* = \log_2 a_i + \log_2 \|v_i - o\| + c$, with $b_i^* = 0$ (eviction) below it. We measure the value term at ≤ 0.035 bits everywhere in our data, so in practice

$$b_i^* = \log_2 a_i + c, \quad (2)$$

which reads attention weights only and never touches the value cache.

Tier extinction and the order parameter. A tier b lies on the lower convex envelope of the cost curve only if storing a token at b bits beats dropping it: $\tau^2 c_b < c_0 = 1$. When $\tau^2 c_b > 1$ the tier is *dead*, meaning no token should ever be stored at b bits. This one inequality produces both edges of the band. At the *sharp* end, rising τ kills rungs from the bottom until only $\{0, b_{\max}\}$ remain, and a two-tier allocation is eviction. At the *diffuse* end, the ladder’s spread, $\text{std}_i(\log_2 a_i) = \tau / \ln 2$, drops below about one rung and the allocation coincides with uniform. We define the **dead-tier fraction** D_2 of a (model, context) configuration as the share of heads with $\tau^2 c_2 > 1$. It uses one calibration pass, has no free parameter, and contains no L .

3 MEASURING THE PHASE DIAGRAM

Setup. We use six models: Llama-3.1-8B and Llama-3.3-70B (Grattafiori et al., 2024), Mistral-7B (Jiang et al., 2023), Qwen1.5-MoE-A2.7B, Qwen3-8B, and Qwen3-30B-A3B (Yang et al., 2025). Each runs at every power-of-two context from 8k up to its native window, giving 24 configurations and 46k head measurements. Prompts are PG-19 haystacks (Rae et al., 2020) with an embedded needle. Every configuration passes an input-validity gate (real text; needle retrieved on every prompt). Keys are quantized with TurboQuant_{mse} (Zandieh et al., 2025) (random rotation plus Lloyd–Max (Lloyd, 1982)), and every reported error is an *exact recomputation* of o with the real quantized keys. Eq. 1 only chooses bit-widths and is never used to score them. At a matched budget of $B=3$ bits/token we compare the water-filling allocation against two corners: uniform B -bit quantization, and eviction keeping $BL/8$ tokens at 8 bits, ranked by H2O accumulated attention (SnapKV/StreamingLLM rankings win $\leq 22\%$ of heads). A head’s *gain* is the best corner’s error divided by the allocation’s error. A head is *in band* if its gain is at least $2\times$, and *routed gain* is the




Figure 2: **(a)** Share of heads where allocation beats the best corner by $\geq 2\times$, versus context (colours as Fig. 1b). Every model decays, at -1.4 to -14.6 points per doubling; the circled point was held out. **(b)** The linearized cost model’s predicted gain over uniform divided by the exactly recomputed gain: all 24 configurations fall within $[0.79, 1.34]$, while the measured gains themselves span $20\text{--}37\times$. **(c)** *Projected, not measured (TBD [R8])*: the hypothesis the router experiment tests. Always-allocate wins at low D_2 and always-evict at high D_2 ; a per-head router should track the upper envelope, with the largest margin in the mixed middle.

geometric mean of $\max(\text{gain}, 1)$, which is what a per-head router can realize. Our primary comparison is *symmetric*: allocation and corner score tokens with the same information. **TBD [R3]: a fully deployable symmetric cell (both sides on lagged attention)** agrees with this one to within 5 band-points on 4 configurations; the re-run with the fixed instrumentation is in flight. Add seeds and error bars on the 128k cells [R7].

Finding 1: the dead-tier fraction is the order parameter. Across 24 configurations D_2 orders routed gain at Spearman $\rho = -0.963$ (Fig. 1b). Configurations nested in six models are not independent points, so we also partial out model identity, which leaves $\rho = -0.944$; within *every* model the relationship is perfectly monotone ($\rho = -1.00$). The obvious alternatives do worse. Logit spread τ alone gives -0.76 pooled, because it does not know where the tiers fall relative to c_0 . Normalized effective support n_{95}/L , a natural “concentration” index, gives $+0.00$, because dividing by L confounds architecture with the context at which it was measured. The order parameter also survives a change of target: replacing the corner’s ranking by a deployable H2O score moves band fractions by 6–29 points yet leaves $\rho = -0.941$. Finally, it passed a held-out test. The boundary was fit without Llama-3.3-70B at 128k, and at that point D_2 jumps from 11.7% to 41.5% and the band falls from 47.5% to 18.1%, as the fitted relation predicts (Fig. 2a, circled).

Finding 2: context length is a phase variable. Within every model τ rises monotonically with L (rank correlation 1.00 in 6/6). Rising τ pushes $\tau^2 c_b$ past c_0 for successively higher rungs, the ladder is eaten from the bottom, and the optimum degenerates toward eviction. In absolute terms allocation keeps improving over uniform with context ($26\times \rightarrow 34\times$ on Llama-3.1-8B); what changes is that the eviction corner catches up. Llama-3.3-70B is the most allocation-friendly model in the study at 8k–64k (78.6% of heads in band at 8k) and among the least at 128k (Fig. 2a). *Architecture rankings are not stable across context*, so a single verdict per model, the way this literature reports results, is not well defined. The same corner can also be reached for a different reason. Qwen3 models sit at 60–73% dead tiers at *every* length because of sink-dominated attention, which is architectural, whereas Llama-3.1-8B reaches the same corner by traversal. **TBD [R4]: attention or RoPE limit? The collapsing Llamas sit at their trained window; Qwen3-30B, at 50% of its window, drops least. Planned: 96k Llamas; Qwen3-30B at 192k/256k.**

Finding 3: the cost model predicts forward. The constant $\tau / \ln 2$ ladder width is an algebraic identity, not evidence. The real test of Eq. 1 is whether it predicts *measured* error ratios. It does: per-configuration median predicted/measured gain lies in $[0.79, 1.34]$ across all 24 configurations (Fig. 2b), for a quantity that spans $20\text{--}37\times$ and was free to be wrong by an order of magnitude.

4 FROM MAP TO METHOD: SIEVE

Any single policy is wrong for some heads at some context, so SIEVE routes per head and per deployment length.

Offline calibration (one forward pass). For each head, measure τ and the per-tier distortion c_b , which gives the set of live tiers and the ladder width. Route each head to *uniform* when its ladder is narrower than about one rung, to *eviction* (H2O) when the tiers below $b_{\max}$ are extinct, and to the *water-filling ladder* otherwise. The ladder’s tier set is read off the convex envelope: tier 0 included, 1-bit excluded. Because τ is monotone in L and D_2 is L -free, a head is

Table 1: Measured facts (24 configurations) and the design decision each one forces.

Measurement	Consequence for SIEVE
Water-filling evicts 40–54% of tokens by itself at $B=3$	Eviction is tier 0 <i>inside</i> the allocator
1-bit tier dead in 95–100% of heads	Never offer 1 bit
Allocation beats uniform 3-bit by 20–37×	The baseline to beat is eviction, not uniform
Oracle token ranking beats H2O by only 1.02–1.41×	Scoring is solved; the value is in <i>allocation</i>
K^*/n_{95} spans 5.8–215, not ordered by support n_{95}	No global keep-count constant; derive per head
Ladder width vs. realized gain: $\rho \in [-0.77, -0.41]$ in 24/24	Cheap per-head routing signal, no quantization

re-phased for a new deployment length by re-evaluating the same inequality, with no new measurement. **TBD [R6]:** *pin the boundary as an interval ($D_2 \in [40, 60]\%$).*

Decode. Band heads run a cascade. A 3-bit base pass over all tokens yields $\hat{a}$; Eq. 2 allocates $b_i^* = \log_2 \hat{a}_i + c$; successively refinable codes (Equitz & Cover, 1991) add +1/ + 2/ + 2 bits on shrinking sets; the tail is evicted. **TBD [R3/R5]:** *re-budgeting cadence. Preliminary: stale scores cost in-band heads 1.4–1.6× but out-of-band heads only 1.00–1.03×, so decode-time re-budgeting may be load-bearing. TBD [R11/R12/R14]: nested-code overhead vs. monolithic 8-bit (kill at ~10%); GQA union cost; kernel TPOT vs. dense FP8 at 128k.*

End-task validation. **TBD [R8]:** Router on versus off on RULER (Hsieh et al., 2024), at matched bits, on the highest- and lowest- D_2 models (Qwen3-30B-A3B; Llama-3.3-70B at $\leq 64k$). The falsifiable prediction (Fig. 2c) is that the router recovers accuracy *exactly where the diagram says the allocator degenerates*, and that the end-task gap shrinks with context as in Fig. 1c. **TBD [R9]:** *K^* -derived per-head budgets vs. Feng et al. (2024); Cai et al. (2024) on LongBench (stands alone).*

5 RELATED WORK AND OUTLOOK

KV quantization ranges from per-channel/per-token integer schemes (Liu et al., 2024; Hooper et al., 2024) to rotation-based and sketching quantizers (Zandieh et al., 2025; 2024). **Mixed precision** (Yang et al., 2024; He et al., 2024) gives salient tokens heuristic higher tiers. **Eviction** scores tokens by accumulated or windowed attention (Zhang et al., 2023; Li et al., 2024; Liu et al., 2023), keeps sinks (Xiao et al., 2024), or reallocates budget across layers and heads (Cai et al., 2024; Feng et al., 2024). Each optimizes *within* a family; we derive the boundary *between* families. Tier extinction also explains why QJL’s unbiased 1-bit tier is never worth offering: it costs more than eviction. We study the *token* axis of the key cache; per-channel key treatment is orthogonal and composes with ours. **TBD [R15]:** *related-work sweep (last 12 months) for concurrent output-distortion and rate-allocation formulations; position against them explicitly.*

Limitations. Every number here is an output-error ratio; whether a 1.3–3× routed gain survives to end-task accuracy and throughput is the main open question (**TBD [R8], [R14]**), and reach beyond 128k hinges on [R4]. The diagram does not depend on either: it says where to spend bits before deployment.

REFERENCES

- Zefan Cai, Yichi Zhang, Bofei Gao, Yuliang Liu, Tianyu Liu, Keming Lu, Wayne Xiong, Yue Dong, Baobao Chang, Junjie Hu, and Wen Xiao. PyramidKV: Dynamic KV cache compression based on pyramidal information funneling. *arXiv preprint arXiv:2406.02069*, 2024.
- Thomas M. Cover and Joy A. Thomas. *Elements of Information Theory*. Wiley, 2nd edition, 2006.
- William H. R. Equitz and Thomas M. Cover. Successive refinement of information. *IEEE Transactions on Information Theory*, 37(2):269–275, 1991.
- Yuan Feng, Junlin Lv, Yukun Cao, Xike Xie, and S. Kevin Zhou. Ada-KV: Optimizing KV cache eviction by adaptive budget allocation for efficient LLM inference. *arXiv preprint arXiv:2407.11550*, 2024.
- Aaron Grattafiori et al. The Llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.
- Yefei He, Luoming Zhang, Weijia Wu, Jing Liu, Hong Zhou, and Bohan Zhuang. ZipCache: Accurate and efficient KV cache quantization with salient token identification. In *Advances in Neural Information Processing Systems*, 2024.
- Coleman Hooper, Sehoon Kim, Hiva Mohammadzadeh, Michael W. Mahoney, Yakun Sophia Shao, Kurt Keutzer, and Amir Gholami. KVQuant: Towards 10 million context length LLM inference with KV cache quantization. In *Advances in Neural Information Processing Systems*, 2024.

- Cheng-Ping Hsieh, Simeng Sun, Samuel Krman, Shantanu Acharya, Dima Rekesh, Fei Jia, and Boris Ginsburg. RULER: What’s the real context size of your long-context language models? In *Conference on Language Modeling*, 2024.
- Albert Q. Jiang et al. Mistral 7B. *arXiv preprint arXiv:2310.06825*, 2023.
- Yuhong Li, Yingbing Huang, Bowen Yang, Bharat Venkitesh, Acyr Locatelli, Hanchen Ye, Tianle Cai, Patrick Lewis, and Deming Chen. SnapKV: LLM knows what you are looking for before generation. In *Advances in Neural Information Processing Systems*, 2024.
- Zichang Liu, Aditya Desai, Fangshuo Liao, Weitao Wang, Victor Xie, Zhaozhao Xu, Anastasios Kyrillidis, and Anshumali Shrivastava. Scissorhands: Exploiting the persistence of importance hypothesis for LLM KV cache compression at test time. In *Advances in Neural Information Processing Systems*, 2023.
- Zirui Liu, Jiayi Yuan, Hongye Jin, Shaochen Zhong, Zhaozhao Xu, Vladimir Braverman, Beidi Chen, and Xia Hu. KIVI: A tuning-free asymmetric 2bit quantization for KV cache. In *International Conference on Machine Learning*, 2024.
- Stuart Lloyd. Least squares quantization in PCM. *IEEE Transactions on Information Theory*, 28(2):129–137, 1982.
- Jack W. Rae, Anna Potapenko, Siddhant M. Jayakumar, and Timothy P. Lillicrap. Compressive transformers for long-range sequence modelling. In *International Conference on Learning Representations*, 2020.
- Guangxuan Xiao, Yuandong Tian, Beidi Chen, Song Han, and Mike Lewis. Efficient streaming language models with attention sinks. In *International Conference on Learning Representations*, 2024.
- An Yang et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.
- June Yong Yang, Byeongwook Kim, Jeongin Bae, Beomseok Kwon, Gunho Park, Eunho Yang, Se Jung Kwon, and Dongsoo Lee. No token left behind: Reliable KV cache compression via importance-aware mixed precision quantization. *arXiv preprint arXiv:2402.18096*, 2024.
- Amir Zandieh, Majid Daliri, and Insu Han. QJL: 1-bit quantized JL transform for KV cache quantization with zero overhead. *arXiv preprint arXiv:2406.03482*, 2024.
- Amir Zandieh, Majid Daliri, Majid Hadian, and Vahab Mirrokni. TurboQuant: Online vector quantization with near-optimal distortion rate. *arXiv preprint arXiv:2504.19874*, 2025.
- Zhenyu Zhang, Ying Sheng, Tianyi Zhou, Tianlong Chen, Lianmin Zheng, Ruisi Cai, Zhao Song, Yuandong Tian, Christopher Ré, Clark Barrett, Zhangyang Wang, and Beidi Chen. H2O: Heavy-hitter oracle for efficient generative inference of large language models. In *Advances in Neural Information Processing Systems*, 2023.

TBD A.1 Measurement integrity: the validity gates (real-corpus enforcement, needle retrieval, probe/chunking equivalence) and the defects they catch. A.2 The $\tau/\ln 2$ ladder identity as a consistency check; the 2-bit quantizer acts as a universal logit temperature ($\alpha = 0.941\text{--}0.944$). A.3 Findings and negative results: why τ and n_{95}/L fail as regime variables; the refuted absolute-support eviction budget; H2 (2-bit score ranking, $\rho = 0.52\text{--}0.77$) and the move to a 3-bit base layer; Gaussian-logit simulation as a qualitatively wrong regime map. A.4 Per-model tables at every context.

CONTENTS

1	Introduction	1
2	One objective: eviction is quantization at zero rate	2
3	Measuring the phase diagram	2
4	From map to method: SIEVE	3
5	Related work and outlook	4

A Introduction: The KV-Cache Compression Dilemma	6
B Output Distortion and the Equivalence Principle	7
B.1 Standard Attention Formulation	7
B.2 Quantization Perturbation Model	7
B.3 First-Order Taylor Approximation of Output Error	7
B.4 Eviction as Zero-Bit Quantization	7
C Optimal Rate Allocation via Reverse Water-Filling	8
C.1 The Constrained Optimization Problem	8
C.2 Lagrangian Solution	8
C.3 The Value-Free Simplification	8
D Tier Extinction and the Order Parameter	8
D.1 The Convexity Condition	8
D.2 The Two Asymptotic Regimes	8
D.3 The Dead-Tier Fraction (D_2)	9
E Empirical Evaluation Methodology	9
E.1 Experimental Architecture	9
E.2 Quantization Engine and Scoring Integrity	9
E.3 Matched-Budget Comparative Baselines ($B = 3$ bits/token)	9
E.4 Evaluation Metrics	9
F Summary of Key Takeaways	9

A INTRODUCTION: THE KV-CACHE COMPRESSION DILEMMA

In autoregressive Transformer generation, storing key and value vectors across long context windows consumes substantial GPU high-bandwidth memory (HBM). Prior research has largely diverged into two competing paradigms:

1. **Token Eviction (Sparsification):** Discarding tokens perceived to be uninformative (e.g., StreamingLLM, H₂O, SnapKV) while retaining retained tokens at full precision.
2. **Uniform Quantization:** Lowering the numerical bit-width of all cached key and value vectors uniformly across the context (e.g., 3-bit or 4-bit uniform quantization).

The theory presented here unifies these two disparate strategies under a single information-theoretic rate-distortion framework, proving that token eviction is not an ad-hoc heuristic, but rather the boundary case of zero-bit quantization.

B OUTPUT DISTORTION AND THE EQUIVALENCE PRINCIPLE

B.1 STANDARD ATTENTION FORMULATION

For a query vector $q \in \mathbb{R}^d$, key vectors $\{k_i\}_{i=1}^L \subset \mathbb{R}^d$, and value vectors $\{v_i\}_{i=1}^L \subset \mathbb{R}^{d_v}$:

$$s_i = \frac{q^\top k_i}{\sqrt{d}} \quad (\text{Pre-softmax logits}), \quad (3)$$

$$a_i = \frac{\exp(s_i)}{\sum_j \exp(s_j)} = [\text{softmax}(s)]_i \quad (\text{Attention weights}), \quad (4)$$

$$o = \sum_{i=1}^L a_i v_i \quad (\text{Attention head output}). \quad (5)$$

B.2 QUANTIZATION PERTURBATION MODEL

When key k_i is quantized using b_i bits, the resulting logit s_i undergoes a perturbation δ_i :

$$\tilde{s}_i = s_i + \delta_i, \quad \text{Var}(\delta_i) = \tau^2 c_{b_i}, \quad (6)$$

where:

- τ denotes the **head logit spread** (the dynamic range or standard deviation of the logits for that specific attention head).
- c_{b_i} is the **relative distortion factor** of the quantizer at bit-width b_i . For standard quantizers, distortion decays exponentially: $c_b \propto 4^{-b} = 2^{-2b}$.
- Any constant bias in the perturbation is completely neutralized because the softmax operator is shift-invariant: $\text{softmax}(s + C \cdot \mathbf{1}) = \text{softmax}(s)$. Thus, mean shift costs nothing, and only the variance $\text{Var}(\delta_i)$ induces output error.

B.3 FIRST-ORDER TAYLOR APPROXIMATION OF OUTPUT ERROR

Differentiating the output vector o with respect to logit s_i :

$$\frac{\partial o}{\partial s_i} = \sum_j \frac{\partial a_j}{\partial s_i} v_j = a_i (v_i - o). \quad (7)$$

Assuming independent perturbations across tokens, the expected squared distortion in the head output vector o to first order is:

$$\mathbb{E} \|\Delta o\|^2 \approx \sum_{i=1}^L a_i^2 \|v_i - o\|^2 \tau^2 c_{b_i}. \quad (8)$$

B.4 EVICTION AS ZERO-BIT QUANTIZATION

If a subset of tokens E is evicted entirely and the remaining attention weights are renormalized, the leading-order output perturbation is:

$$\|\Delta o_{\text{evict}}\|^2 \approx \sum_{i \in E} a_i^2 \|v_i - o\|^2. \quad (9)$$

Comparing this directly to Equation equation 8, eviction precisely matches the quantization error formula when:

$$\tau^2 c_0 = 1 \implies c_0 = \frac{1}{\tau^2}. \quad (10)$$

Remark B.1 (Fundamental Result). Eviction is identical to a quantizer rung at $b_i = 0$. Its cost constant c_0 is mathematically derived from the curvature of the softmax operator rather than being hand-tuned as an empirical penalty.

C OPTIMAL RATE ALLOCATION VIA REVERSE WATER-FILLING

C.1 THE CONSTRAINED OPTIMIZATION PROBLEM

Given a total bit budget of BL bits across L tokens (an average precision of B bits per token), the objective is to choose $\{b_i\}_{i=1}^L$ to minimize output distortion:

$$\min_{\{b_i\}} \sum_{i=1}^L a_i^2 \|v_i - o\|^2 \tau^2 c_{b_i} \quad \text{subject to} \quad \sum_{i=1}^L b_i \leq BL, \quad b_i \geq 0. \quad (11)$$

C.2 LAGRANGIAN SOLUTION

Using the distortion model $c_{b_i} \propto 4^{-b_i} = 2^{-2b_i}$, the Lagrangian separates independently for each token i . Setting the gradient with respect to b_i to zero yields the water-filling solution:

$$b_i^* = \log_2 a_i + \log_2 \|v_i - o\| + c, \quad (12)$$

where c is a normalization constant determined by the total bit budget BL .

- Tokens whose allocated bits exceed the threshold (above the “water level”) are assigned b_i^* .
- Tokens falling below the threshold are dropped to $b_i^* = 0$ (**eviction**).

C.3 THE VALUE-FREE SIMPLIFICATION

Empirical measurements show that the variation of the value distance term $\log_2 \|v_i - o\|$ is negligible across tokens (≤ 0.035 bits throughout real-world context data). Consequently, the value term can be absorbed into the constant:

$$b_i^* \approx \log_2 a_i + c. \quad (13)$$

Operational Significance: Optimal bit allocation can be computed entirely from the attention weights a_i alone, without requiring memory access to the value cache.

D TIER EXTINCTION AND THE ORDER PARAMETER

D.1 THE CONVEXITY CONDITION

On real hardware accelerators, bit-widths cannot vary continuously; they must belong to discrete rungs (e.g., $b \in \{0, 2, 4, 8\}$). A discrete tier b is viable (lies on the lower convex envelope of the rate-distortion curve) if and only if storing a token at b bits introduces less distortion than evicting it:

$$\tau^2 c_b < \tau^2 c_0 = 1. \quad (14)$$

Definition D.1 (Dead Tier). If $\tau^2 c_b > 1$, tier b is defined as **dead**. Storing a token at b bits injects more variance into the logit than completely deleting the token. Therefore, no token should ever be quantized to b bits.

D.2 THE TWO ASYMPTOTIC REGIMES

The inequality $\tau^2 c_b < 1$ determines the compression topology based on the head’s logit spread τ :

1. **The Sharp End (High τ , Large Dynamic Range):** As τ increases, $\tau^2 c_b$ exceeds 1 for low-bit precisions, killing rungs from the bottom up (e.g., 2-bit dies first, then 3-bit). Eventually, the available tier set collapses strictly to $\{0, b_{\max}\}$.

Optimal strategy $\longrightarrow$ **Pure Eviction (Keep at full precision or drop).**

2. **The Diffuse End (Low τ , Small Dynamic Range):** When τ is small, the standard deviation of the log-attention weights collapses:

$$\text{std}_i(\log_2 a_i) = \frac{\tau}{\ln 2} \rightarrow 0. \quad (15)$$

The ladder spread shrinks below the width of a single rung, meaning all tokens receive essentially the same bit allocation.

Optimal strategy $\longrightarrow$ **Uniform Quantization.**

D.3 THE DEAD-TIER FRACTION (D_2)

To quantify model behavior, the **dead-tier fraction** D_2 is defined as:

$$D_2 = \frac{1}{H} \sum_{h=1}^H \mathbb{I}(\tau_h^2 c_2 > 1), \quad (16)$$

representing the proportion of attention heads where 2-bit quantization is inferior to eviction. D_2 acts as an order parameter: it requires only a single calibration pass, possesses zero free parameters, and is invariant to sequence length L .

E EMPIRICAL EVALUATION METHODOLOGY

E.1 EXPERIMENTAL ARCHITECTURE

To test the analytical theory on real LLMs, an empirical testbed is structured as follows:

- **Scale:** 24 distinct model/context configurations evaluated across 46,000 individual attention heads.
- **Context Scaling:** Context windows span every power-of-two from 8k tokens up to native model context limits.
- **Workload:** Synthetic “needle-in-a-haystack” retrieval embedded inside natural PG-19 book text.
- **Input-Validity Gate:** Strict verification ensuring that the uncompressed model successfully retrieves the needle on every prompt before testing compression.

E.2 QUANTIZATION ENGINE AND SCORING INTEGRITY

- **Quantizer:** Key vectors are quantized using TurboQuant_{mse}, which combines random orthogonal rotations with Lloyd–Max optimal scalar quantization.
- **Strict Separation:** Equation equation 8 and Equation equation 13 are utilized *exclusively* to plan bit assignments (b_i^*). The recorded error is always the **exact recomputed output** o evaluated through the real quantized key vectors.

E.3 MATCHED-BUDGET COMPARATIVE BASELINES ($B = 3$ BITS/TOKEN)

All techniques operate under an identical bit budget ($B = 3$ bits per token):

Compression Scheme	Mechanism	Regime Alignment
Uniform Quantization	Every token stored at exactly $b_i = 3$ bits	Diffuse Corner
Pure Eviction	Keep top $BL/8 = 3L/8$ tokens at 8 bits; evict the rest	Sharp Corner
Water-Filling	Allocates $b_i^* \approx \log_2 a_i + c$ across discrete tiers	Intermediate Spectrum

Table 2: Matched-budget comparison at 3 bits/token. Eviction ranks tokens via H₂O accumulated attention.

E.4 EVALUATION METRICS

- **Per-Head Gain:** The ratio of the best corner baseline error to the water-filling error:

$$\text{Gain} = \frac{\min(\text{Error}_{\text{uniform}}, \text{Error}_{\text{eviction}})}{\text{Error}_{\text{water-filling}}}. \quad (17)$$

- **In-Band Heads:** Attention heads exhibiting $\text{Gain} \geq 2\times$, where mixed-precision water-filling outperforms both classical extremes by at least a factor of two.
- **Routed Gain:** The geometric mean of $\max(\text{Gain}, 1)$, reflecting the real-world performance attained by a dynamic per-head router selecting the optimal compression scheme.
- **Symmetric Information Principle:** Allocation and corner baselines score token importance using identical attention statistics, ensuring fair benchmarking.

F SUMMARY OF KEY TAKEAWAYS

Core Concept	Key Takeaway
Eviction is 0-bit quantization	Setting $\tau^2 c_0 = 1$ naturally aligns token dropping with quantization theory; no separate eviction heuristic is needed.
Logit spread τ dictates policy	High τ (sharp heads) kills low-bit tiers and favors pure eviction; low τ (diffuse heads) favors uniform quantization.
Value-cache independence	Bit allocation depends almost entirely on $\log_2 a_i$, enabling optimal allocation without accessing or loading value vectors.
Hardware routing potential	Modeling in-band heads enables per-head routing, optimizing accuracy per byte across heterogeneous attention layers.

Table 3: Summary of foundational insights.